

Position: Web Agent Research Should Prioritize Capture-Replay Over Handcrafted Replicas

Anonymous Authors¹

Abstract

This is a position paper. Building realistic web agent environments is extraordinarily expensive; TheAgentCompany required 3,000 person-hours for 175 tasks. This cost barrier has spawned a proprietary ecosystem where well-funded labs train on environments they cannot share, splitting the field into haves and have-nots. We argue this trajectory is unsustainable: the community should prioritize capture-replay infrastructure over handcrafted website replicas. A single expert demonstration, recorded on a live website, can produce a complete offline environment in minutes. We present TRACE, a proof-of-concept that captures production sites including authenticated flows, and validate scalability: 6 fully replayable environments captured in 40 minutes (~ 7 min/task), plus 71 tasks collected in a pilot study, representing a $120\times$ reduction in marginal time compared to replica construction. Capture-replay has real limitations; agents that deviate from recorded trajectories encounter missing responses, and human demonstrations cannot anticipate every valid agent strategy. But these coverage gaps are explicit and measurable (agents hit missing responses), whereas replica fidelity gaps (where simulated behavior diverges from production) may go undetected. We call on the community to shift investment toward capture tooling, large-scale collection, and open sharing. The tools exist; the decision is whether we continue paying the replica tax or move the field forward.

1. Introduction

The past two years have witnessed remarkable progress in language-model agents for web and computer use. Systems

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

built on frontier models can now navigate complex websites, fill forms, execute multi-step workflows, and retrieve information across diverse domains. Benchmarks such as Mind2Web (Zhang et al., 2023), WebArena (Zhou et al., 2023), WebVoyager (He et al., 2024), WorkArena (Drouin et al., 2024), REAL (Garg et al., 2025), OSWorld (Xie et al., 2024), BEARCUBS (Song et al., 2025), BrowseComp (Wei et al., 2025), and TheAgentCompany (Xu et al., 2024) have provided standardized evaluation environments that enable reproducible comparisons and have driven rapid capability improvements.

Yet beneath this progress lies a troubling structural problem: **the environments we use to train and evaluate web agents bear little resemblance to the tasks that would make these agents economically valuable.** The majority of benchmark tasks fall into categories we characterize as “deep research” (trivia-style multi-hop questions that rarely require browser automation), “information seeking” (simple lookups that could often be answered by a search engine), or “atomic execution” (short, isolated actions like clicking a button or filling a single field). Long-horizon, multi-step workflows that mirror real paid work (booking complex travel itineraries, configuring enterprise SaaS tools, executing financial transactions, managing e-commerce operations) remain dramatically underrepresented. More importantly, most benchmarks are not collected from people doing paid work; they are task templates designed for evaluability rather than traces of real economic activity.

This is not an oversight. It reflects a fundamental constraint: **handcrafted environments are extraordinarily expensive to build.**

Our position: The web agent research community should prioritize capture-replay infrastructure over handcrafted replicas. Rather than continuing to invest primarily in website replicas, we should record experts completing tasks on live websites and transform those recordings into reusable, offline environments. A single expert demonstration can produce a complete environment in minutes, compared to the hours or days required to handcraft a replica. Replicas retain value for controlled ablations, but the marginal return on capture-replay investment far exceeds that of additional

Table 1. Environment construction costs for prominent benchmarks.

Benchmark	Tasks	Reported Effort
WebArena	~812	Months of development; self-hosting four functional website replicas
OSWorld	~369	Hand-annotated tasks with initial state scripting and custom evaluators
TheAgentCompany	~175	3,000 person-hours by 21 contributors over two months
REAL	~112	Deterministic replicas of 11 real websites with custom harnesses
TRACE (ours)	6[†]	300.7 hours: 300 tooling + 0.7 collection (~7 min/task)

[†]Fully replayable; 71 additional tasks collected in pilot (Sec. 3.3).

replica construction.

Consider the costs documented in recent benchmark papers:

TheAgentCompany’s transparent reporting is particularly illuminating: building 175 realistic “employee-style” tasks required 3,000 person-hours, the equivalent of 1.5 full-time engineers working for an entire year. The project involved 21 contributors including software engineers and project managers, with complex tasks requiring more than 10 hours each to design, implement, test, and verify. And these are among the most resource-rich academic teams in the field.

The implication is stark: at current construction costs, the academic community cannot build environments at the scale needed to cover the long tail of economically valuable browser tasks. We can produce hundreds of tasks, perhaps low thousands with extraordinary effort, but the space of valuable web workflows spans millions of distinct patterns across hundreds of thousands of websites.

These costs are not only financial. They consume scarce research time that could otherwise go toward agent design, failure analysis, and iteration. And because realistic environment construction is so expensive, evaluation suites remain small and often collapse to binary success metrics, which limits error analysis and makes it difficult to diagnose where, why, and how models fail.

This gap has not gone unnoticed by industry. A growing ecosystem of startups now builds browser environments for AI agent training (mec, 2025; pla, 2025; agi, 2025; kai, 2025). The business models vary: some host live environments, others deliver static datasets, and some offer both. What they share is a focus on creating replica websites and simulated workflows where AI agents can learn through reinforcement learning without triggering blocking mechanisms on real sites.

The scale of investment is remarkable. According to *The Information*, Anthropic has discussed spending more than \$1

billion on RL environments over the next year (Zeff, 2025). Data-labeling giants like Scale AI, Surge, and Mercor are pivoting resources toward environment construction, with Surge reportedly generating \$1.2 billion in revenue last year from AI lab contracts and recently spinning up a dedicated internal organization for RL environments. Mechanize, a startup focused exclusively on environments, offers software engineers \$500,000 salaries to build them, far exceeding typical data-labeling contractor rates (Zeff, 2025).

Leading AI labs treat these environments as core training infrastructure. **Browser environments have become a strategic asset: valuable, proprietary, and concentrated in the hands of well-funded organizations.**

The field is at an inflection point. Three trends are converging: (1) agents are becoming capable enough that evaluation on short, synthetic tasks no longer discriminates between systems, creating demand for long-horizon, realistic environments; (2) the cost of building such environments has created a two-tier ecosystem where frontier labs invest billions while academic groups work with hundreds of tasks; and (3) the resulting concentration raises safety concerns, as independent researchers cannot study systems trained on environments they cannot access. If the community does not change course, the gap will widen and become structural.

We believe this trajectory is problematic for the field. When the ability to train and iterate on realistic web agents depends on access to expensive proprietary infrastructure, the research community’s capacity for independent investigation is constrained. Open science suffers. Reproducibility suffers. And the concentration of capability in a small number of actors raises broader concerns about the development trajectory of increasingly powerful autonomous systems. This issue affects not only academic groups, but also open-source LM companies operating on budgets that are tiny relative to frontier labs.

This paper argues for an alternative approach: capture-replay.

The core idea is simple: rather than handcrafting website replicas, we record an expert completing a task on a live website and transform that recording into a self-contained environment that can be replayed offline. A single expert demonstration (taking minutes to perform) produces a complete environment bundle including DOM states, network responses, screenshots, and interaction logs. Agents can then interact with this frozen snapshot without contacting the live web, with coverage expanding as additional trajectories are captured.

Capture-replay is not a new concept in software engineering. Record-and-replay debugging, network mocking, and browser automation testing have used similar techniques for decades. But its application to web agent research has been

surprisingly limited. We argue this represents a significant missed opportunity.

To demonstrate feasibility, we developed TRACE (Trajectory Recording and Capture Environments), an open-source capture-replay pipeline. Building robust tooling required approximately 300 hours of engineering effort: analyzing hundreds of captured sessions, developing URL canonicalization heuristics, and iterating on replay matching. But once built, the tooling amortizes: we captured six diverse environments (GitHub, Amazon, Airbnb, Kayak, Uniqlo, Ultimate Guitar) in approximately 40 minutes of total collection time, including retries and quality checks. A separate pilot study collected 71 additional tasks in 10 hours. This is roughly 7 minutes per task versus the 17+ hours per task reported for handcrafted benchmarks like TheAgentCompany, representing a **120× reduction in marginal collection time**.

We emphasize “prioritize” deliberately. Handcrafted replicas have absorbed enormous research effort while capture-replay infrastructure remains underexplored. Replicas offer value for controlled ablations where researchers need to systematically vary specific elements, but determinism, breadth, and realistic task coverage can also be achieved by capture-replay at scale through multi-trajectory collection, record-on-miss expansion, and post-processing that enables controlled perturbations. We have not yet demonstrated these techniques; our proof-of-concept validates single-trajectory capture and replay. But the economics favor building this infrastructure: even if adequate coverage requires 10 trajectories per task, that is still an order of magnitude cheaper than replica construction.

2. The Current Landscape: Expensive Replicas, Limited Coverage

2.1. A Taxonomy of Browser Agent Benchmarks

To understand the current state of web agent environments, we systematically analyzed eleven prominent benchmarks across three dimensions: **task category** (what kind of work the agent performs), **task horizon** (how many steps or how much time tasks typically require), and **economic grounding** (whether tasks resemble work someone would pay to have completed).

Several patterns emerge from this analysis. First, an **effort-value tradeoff**: benchmarks emphasizing economically realistic tasks (TheAgentCompany, WorkArena, REAL) require dramatically more construction effort than those emphasizing information retrieval or synthetic reasoning (BrowseComp, GAIA). Realistic execution tasks require functional environments with state that actually changes, authentication flows that work, and evaluation criteria that verify real outcomes. Meanwhile, existing benchmarks heavily favor a

small set of domains (e-commerce, developer tools, travel), leaving vast categories of economically important web work (financial services, healthcare, government, enterprise SaaS) largely uncovered.

Second, **horizon limitations**: even benchmarks labeled “long-horizon” often consist of relatively short interaction sequences (Li et al., 2025), and as the field pushes toward longer-horizon evaluations (METR, 2025), handcrafting scales poorly. If some tasks already require 10+ hours to design, multi-day workflows are untenable. Replica construction cannot keep pace with the horizon lengths we need to measure.

Third, a **live-web evaluation problem**: benchmarks on live websites (Zhang et al., 2023; Wei et al., 2025; Song et al., 2025; He et al., 2024) face continuous validity challenges as the web changes, while replica-based benchmarks (Zhou et al., 2023; Garg et al., 2025) gain reproducibility but lose coverage. Sensitive operations (payments, authentication) cannot be retried freely without real consequences.

2.2. The Cost Structure of Environment Construction

Why are realistic browser environments so expensive? The costs compound across multiple dimensions. **Construction costs** include reverse-engineering website interaction models, implementing backend logic, and scripting state initialization (OSWorld required extensive scripting for 370 tasks (Xie et al., 2024); TheAgentCompany built an entire simulated company (Xu et al., 2024)). **Maintenance costs** accumulate as the web changes constantly; both live-site benchmarks and replicas must track upstream changes to remain valid. **Evaluation costs** arise when tasks have multiple valid solutions, pushing the field toward LLM-as-a-judge approaches with their own biases and reproducibility concerns (Xue et al., 2025). Finally, **credential costs** force benchmarks to either avoid authenticated tasks, use synthetic accounts with limited functionality, or require evaluators to provide their own credentials, each constraining coverage or reproducibility.

2.3. The Concentration of Environment Access

The expense of environment construction has predictable consequences for who can build and access high-quality environments.

As discussed above, a commercial ecosystem has emerged to fill this gap. These providers offer browser environments through various models: some host live replicas, others deliver packaged datasets, and some provide both. The environments enable something impossible on real websites: running thousands of AI agents simultaneously through trial-and-error learning without being blocked.

Leading AI labs use these commercial environments for

Table 2. Taxonomy of existing web agent benchmarks. We observe a clear pattern: benchmarks with high economic grounding (TheAgentCompany, WorkArena, parts of REAL) require the most construction effort, while scalable benchmarks (Mind2Web, BrowseComp) tend toward tasks with limited economic value.

Benchmark	Category	Horizon	Econ. Grounding	Key Limitation
GAIA	Deep research	Medium	Low	General-assistant QA; many tasks not web-specific
BEARCUBS	Info seeking	Short/Medium	Low	QA-only; no state-changing actions; small (111 tasks)
BrowseComp	Deep research	Long	Low	QA-only; short answers; open-web drift
Mind2Web 2	Deep research	Long	Low	Agentic search QA; LLM-as-judge evaluation
Mind2Web	Info seeking / Exec	Short/Medium	Medium	Live-web trajectories; action-matching eval; no deterministic env
WebVoyager	Info seeking / Exec	Medium	Low/Medium	Live websites; GPT-4V-based eval; limited domain coverage
WebArena	Execution	Medium	Medium	Replica sites; limited domains; high construction cost
REAL	Execution	Short/Medium	Medium/High	Deterministic replicas; limited task count and domains
OSWorld	Execution	Short/Medium	Medium	Desktop-focused; heavy state setup and evaluators
WorkArena	Execution	Short/Medium	High	Single platform (ServiceNow); small (33 tasks); remote-hosted
TheAgentCompany	Execution	Medium	High	Realistic workflows but extraordinarily expensive to build

training and evaluating their web agents. As frontier labs have exhausted most available text data for pretraining, reinforcement learning in simulated environments has become increasingly important. With per-task costs in the thousands to tens of thousands of dollars, these environments represent substantial investments, but ones that well-capitalized organizations can afford while academic groups generally cannot.

We also see this dynamic inside product companies building agents for their own products: teams are spun up to build replicas of their own UIs and workflows so they can train and evaluate at scale. This is yet another signal that handcrafted environments are expensive, bespoke infrastructure, not a sustainable research path.

These training environments are typically proprietary and contract-bound, making it hard for open-source and academic research to access or reproduce the training conditions that drive frontier results.

The result is a two-tier system: well-funded labs train agents on high-quality, realistic environments that they cannot share, while the academic community and open-source LM companies work primarily with the limited set of open benchmarks. This concentration raises concerns across multiple dimensions. First, *reproducibility* suffers: if state-of-the-art results depend on proprietary training environments, the research community cannot fully verify or build upon them. Second, *diversity* may narrow: concentrated environment access can lead to agents that excel on specific

task distributions while failing on the broader diversity of real-world needs. Third, *safety research* is constrained: independent safety research requires access to the same environments used for capability development, and concentration of access limits this critical work.

3. The Case for Capture-Replay

3.1. Core Concept

Capture-replay inverts the environment construction process. Rather than building a website replica and then defining tasks on it, we follow a three-stage pipeline. In the **capture** stage, an expert performs a real task on a live website while an instrumented browser records everything: DOM states, network traffic, user interactions, screenshots, and video. Next, the **processing** stage transforms raw captures into clean, reusable artifacts: high-level action sequences (a standardized tool-call DSL), extracted credentials (for secure handling), semantic checkpoints (for partial-credit evaluation), and filtered network logs (removing analytics and non-essential traffic). Finally, in the **replay** stage, a replay engine serves the captured assets locally, allowing agents to interact with a frozen snapshot of the original session without contacting the live web.

Relation to prior record-replay systems. Record-replay is well-established in software engineering. Syscall-level tools like Mozilla rr (O’Callahan et al., 2017) enable deterministic debugging of C/C++ programs by recording and

replaying system calls, but operate below the HTTP layer and cannot mock web responses. HTTP-level tools such as Polly.js (Netflix, 2018), VCR (Binié, 2011), and Nock (Azevedo, 2011) record network traffic for unit testing, using strict URL-and-method matching that fails when requests contain dynamic parameters (timestamps, session tokens, randomized IDs). Playwright’s HAR replay (Microsoft, 2023) provides basic glob-pattern matching but lacks semantic understanding of request equivalence.

TRACE extends these foundations with a **multi-tier matching system** designed for the non-determinism inherent in web agent evaluation: (1) exact URL matching when possible, (2) character-frequency similarity scoring for URLs with dynamic query parameters, and (3) LLM-based disambiguation when multiple candidates remain. Crucially, TRACE tracks which HAR entries have been consumed to prevent response reuse, which is essential for flows where the same URL is requested multiple times with different expected responses (e.g., login then authenticated requests). These additions address the gap between unit-test mocking, which assumes deterministic request sequences, and agent evaluation, where exploration creates novel request orderings.

This approach offers three structural advantages. First, **scalability**: a single expert demonstration produces a complete environment in the time it takes to perform the task. Our proof-of-concept captured six diverse tasks in about 40 minutes total (6:40 per task), each producing hundreds of HTTP responses and dozens of DOM snapshots that would require days to handcraft. At these speeds, a single researcher could produce hundreds of environments per week, crowdsourced collection becomes feasible, and any site an expert can use becomes a potential environment.

Second, **economic grounding**: captured environments reflect real tasks by construction. When an expert books an actual flight or configures a real SaaS tool, the resulting environment captures a workflow someone demonstrably values. This contrasts with replica-based benchmarks where researchers imagine valuable tasks and then construct environments to support them. Capture-replay inverts this: **find people doing economically valuable work, record them, and build environments from their actual workflows**.

Third, **democratization**: open-source capture-replay tooling breaks the concentration of environment access. Any researcher with a browser can record environments from any website they can access, with no proprietary infrastructure, licensing fees, or vendor lock-in. TRACE is released as open-source software with captured environments published on Hugging Face.¹

¹Code and data will be released publicly upon acceptance.

3.2. TRACE: A Proof of Concept

To demonstrate that capture-replay is technically feasible on modern production websites, we developed TRACE (Trajectory Recording and Capture Environments), an open-source pipeline implementing the full capture-process-replay workflow.

TRACE implements three core components: (1) a **collection** system using a stealth-configured Playwright browser that records DOM states, network traffic, user interactions, screenshots, and browser storage; (2) a **post-processing** pipeline that converts raw events to a standardized action DSL, extracts credentials for secure handling, identifies semantic checkpoints, and filters non-essential traffic; and (3) a **replay** engine that serves captured assets offline with deterministic request matching. Full implementation details are provided in Appendix A.

Demonstration dataset. We release six captured environments spanning GitHub, Amazon, Ultimate Guitar (authenticated flows with login), and Uniqlo, Kayak, Airbnb (unauthenticated). These cover e-commerce, travel search, social coding, and media sites, including sensitive interactions like payment flows. Statistics are provided in Appendix C.

The dataset is intentionally small: six tasks demonstrates *feasibility*, not comprehensive coverage. Building TRACE required approximately 300 hours of engineering effort, but once built, the six environments were captured in about 40 minutes total (~6 to 7 minutes per task). This asymmetry (high tooling cost, low marginal collection cost) is precisely why capture-replay scales where replicas do not.

3.3. Scaling Validation: A 71-Task Pilot Study

To validate that capture-replay scales beyond proof-of-concept, we conducted a pilot study using an earlier version of the TRACE tooling. A single non-expert data collector reproduced 71 tasks from the Mind2Web benchmark in approximately 10 hours of collection time (roughly 8.5 minutes per task including retries and verification).

What the pilot demonstrates. The collection time data validates our core scalability claim independently of replay quality. Even with imperfect tooling, a single collector working at modest pace (\$4/hour) produced 71 task demonstrations for approximately \$40 USD total. This represents a **120× reduction in per-task time** compared to TheAgentCompany’s reported 17 hours per task for handcrafted environments.

What the pilot revealed. The initial captures were not immediately replayable. The pilot used an earlier data format that lacked proper HAR (HTTP Archive) recordings; the network traffic was stored in a different structure that cannot

Table 3. Cost comparison: handcrafted replicas vs. capture-replay. Replicas have high per-task costs; capture-replay has high fixed costs but near-zero marginal costs.

	TheAgentCompany (Replicas)	TRACE (ours) (Capture-Replay)
Tasks created	175	6 (+ 71 pilot)*
Fixed costs (tooling)	included	300 hrs
Collection/build time	3,000 hrs [‡]	0.7 hrs (+ 10 pilot)
Time per task	17.1 hrs	~7 min
Cost per task[†]	\$855	\$0.47
Speedup factor	n/a	120×
Break-even point*	n/a	~35 tasks

*6 fully replayable; 71 pilot tasks collected but require format migration.

[‡]Includes custom programmatic checkpoint evaluators; TRACE uses LLM-as-a-judge.

[†]Replica cost assumes \$50/hr engineering rate; capture-replay uses \$4/hr collection rate.

*Break-even assumes \$30k tooling investment (300 hrs × \$100/hr).

be directly replayed by TRACE’s current architecture. This experience directly informed the tooling improvements that make the six demonstration environments fully functional.

Why this matters. The pilot illustrates a key property of capture-replay that replicas lack: **collection failures are cheap**. When handcrafted replica construction fails, weeks of engineering effort may be lost. When capture fails, recollection costs minutes. The 71-task pilot “failed” in the sense that those specific captures require substantial refactoring to be replayable, but the collection time investment was 10 hours, not 3,000. We can simply recollect with improved tooling.

Cost comparison. Table 3 presents a direct comparison of environment construction costs across methodologies. We use TheAgentCompany as a proxy because they transparently reported costs, a practice we commend and encourage. We acknowledge this comparison has caveats: TheAgentCompany is a desktop/computer-use benchmark (not web-only), and their 3,000 hours includes custom programmatic evaluation functions, not just environment construction. TRACE uses LLM-as-a-judge for checkpoint evaluation, which requires less per-task engineering but may differ in evaluation precision. Despite these differences, the order-of-magnitude gap in marginal collection time (17 hours vs. 7 minutes) reflects a real structural difference between handcrafted and capture-replay approaches.

The 71-task pilot, despite its limitations, provides empirical grounding for the economic argument: **capture-replay marginal costs are two orders of magnitude lower than replica construction**, and this advantage compounds as collection scales.

4. Alternative Views and Counterarguments

A position paper must engage seriously with opposing views. We address the most significant counterarguments to capture-replay:

4.1. “Capture-replay limits agent exploration”

The concern: Captured environments are constrained to recorded trajectories; if an agent deviates, the replay may lack needed responses.

Our response: This concern is legitimate, and we observe it in practice. When agents deviate from recorded trajectories, they encounter requests for which no response exists in the HAR file. This is a real limitation of single-trajectory capture.

For many deviation types, this is a **coverage problem addressable through data collection**. Multi-trajectory collection captures multiple expert paths through the same task, providing recovery behaviors and alternatives. **Record-on-miss expansion** adds missing branches by capturing additional sessions when agents deviate. These techniques can broaden coverage along dimensions where humans naturally vary: different filter orderings, alternative navigation paths, varied form-filling sequences.

A deeper limitation. We acknowledge that the concern extends beyond coverage quantity. Human demonstrations inherently anchor environments around human-legible solution paths. But increasingly capable agents may discover valid strategies that humans do not naturally use: exhaustively enumerating DOM elements, probing URLs directly, exploiting implicit page structure, or pursuing action sequences that seem inefficient to human operators but reliably achieve correct outcomes. These behaviors may be effective yet absent from any reasonable number of human trajectories. A capture-based environment risks penalizing or failing such behaviors simply because they were never demonstrated.

This is a genuine limitation of the capture-replay paradigm that we do not claim to solve. Multi-trajectory collection cannot anticipate fundamentally non-human strategies. However, we argue capture-replay still represents progress: coverage gaps become *measurable* (how often do agents hit missing responses? which request patterns fail?) rather than *invisible* (do replicas faithfully model production site behavior? how would we know?). Replicas face their own fidelity gaps: they are approximations of real sites, and the approximation errors are often uncharacterized. Capture-replay surfaces its failure modes explicitly.

Economic argument. Crucially, the economics still favor capture-replay even accounting for this limitation. If a task

requires 10 trajectories to achieve adequate coverage for human-like strategies, that represents roughly 70 minutes of collection time (10×7 min), compared to 17+ hours for a single handcrafted replica. The $120\times$ cost advantage shrinks but remains substantial. For strategies that no human trajectory captures, both approaches face challenges, but capture-replay at least makes the gap visible. We have not yet demonstrated multi-trajectory collection at scale; this is future work. But the economic argument for investing in capture infrastructure rather than replica construction holds regardless.

4.2. “Legal and ethical concerns around captured content”

The concern: Capturing and redistributing website content raises legal (copyright, ToS) and ethical (privacy, consent) questions.

Our response: These concerns are legitimate. We advocate for credential extraction and substitution, PII redaction, restricted distribution with research-use agreements, and community ethical guidelines (proposed in Appendix D).

Importantly, replica-based benchmarks face analogous concerns. WebArena clones Reddit’s interface; REAL reproduces real website behavior. Capture-replay does not introduce novel legal risk categories; it inherits the same ambiguities that replica builders have navigated for years, but makes provenance explicit, potentially supporting clearer fair-use arguments.

4.3. “Six tasks doesn’t prove scalability”

The concern: Six tasks doesn’t demonstrate benchmark-scale collection or web diversity.

Our response: We separate two claims: *collection scalability* and *replay quality*. Our 71-task pilot study (Section 3.3) demonstrates collection scalability: a single non-expert collector gathered 71 tasks in 10 hours. The six fully-replayable environments demonstrate replay quality: a human following the original trajectory receives correct responses for every request, with LLM disambiguation resolving ambiguous matches deterministically.

Together, these results demonstrate that capture scales to dozens of tasks with minimal effort, that the replay infrastructure correctly serves recorded responses when the trajectory is followed, that challenging cases including authentication, payments, and dynamic content are handled, and that the 300-hour tooling investment built generalizable infrastructure rather than task-specific code.

What we have not demonstrated: Agent evaluation at scale. As discussed in Section 4.1, agents that deviate from

recorded trajectories encounter missing responses. Validating that capture-replay supports open-ended agent evaluation requires multi-trajectory collection, which we propose but have not yet implemented. The current demonstration validates infrastructure correctness via human replay; expanding coverage for agent evaluation is future work.

The pilot’s captures require refactoring to match TRACE’s current HAR-based format, but this illustrates rather than undermines our argument: when capture fails, recollection costs minutes. When replica construction fails, it costs weeks.

4.4. “Handcrafted replicas are necessary for controlled studies”

The concern: Research requiring controlled ablations or systematic variation needs handcrafted environments.

Our response: We disagree. Multi-trajectory collections can provide alternative paths; post-processing can enable controlled perturbations (DOM edits, resource swaps); deterministic replay provides reproducibility. We have not found research questions that strictly require handcrafted replicas.

4.5. “Commercial providers validate the need for replicas”

The concern: Companies charge \$3,000 to \$50,000 per task. If capture-replay sufficed, why pay these prices?

Our response: Commercial providers *do* use capture-replay methods. They sell engineering labor, not fundamentally different methodology. TRACE required ~ 300 hours of tedious engineering (URL canonicalization, replay debugging, detection logic). Providers monetize this work; our argument is that it should be done once and shared openly.

4.6. “Capture-replay works for evaluation but not training”

The concern: RL training requires exploration and feedback from mistakes that captured environments cannot provide.

Our response: We see no fundamental barrier, only data scale challenges. Multi-trajectory capture could provide a state-action graph suitable for offline RL; checkpoints enable graded rewards; expert demonstrations support imitation learning. The offline RL literature demonstrates learning from fixed datasets is viable.

Evaluation is the more urgent bottleneck. Most academic groups cannot evaluate on realistic long-horizon tasks be-

cause environments do not exist. Capture-replay addresses this immediately.

4.7. Technical Limitations

We acknowledge several technical limitations:

Single-trajectory coverage. The current TRACE demonstration uses single-trajectory captures, validated via human traversal. We have not validated agent evaluation, where deviation from recorded paths causes missing responses. Supporting agent exploration requires multi-trajectory collection, which we propose but have not implemented. This is the most significant gap between our proof-of-concept and production-ready infrastructure.

LLM-based matching non-determinism. TRACE uses LLM disambiguation when multiple responses match a request. This is a proof-of-concept convenience; production systems should default to deterministic matching via URL canonicalization and strict tie-breakers.

Anti-automation resistance. Websites deploy CAPTCHAs, bot detection, and fingerprinting. We found many high-value sites capturable with sufficient engineering, and replay is easier than capture since offline replay avoids triggering anti-bot measures entirely. This is an ongoing arms race, but a practical obstacle rather than a fundamental methodological barrier.

Temporal validity. Captures are snapshots; replaying a task after capture may encounter mismatches as dynamic content diverges. Date-sensitive tasks (e.g., travel search) degrade faster than static workflows (e.g., authentication); see Appendix C for validation data showing this variance. Practical mitigations include periodic recapture and date-aware URL normalization. Replica-based benchmarks face analogous drift but hide it; capture-replay makes temporal validity explicit and measurable.

5. Call to Action

We call on the community to make four shifts:

Prioritize capture tooling over replica construction. The marginal return on capture-replay investment far exceeds that of additional handcrafted replicas. Redirect engineering effort toward capture-replay infrastructure that can produce environments 120× faster. Research groups with domain expertise (finance, healthcare, enterprise SaaS) should capture environments in their areas rather than waiting for generalist benchmarks. Replicas retain value for controlled ablations, but should not be the default methodology.

Document construction costs transparently. The Agent-

Company’s disclosure of 3,000 person-hours set a valuable precedent. All benchmark papers should report environment construction effort to enable informed methodology comparisons.

Share captured environments openly. Publish environments under research-use licenses with credential substitution and PII redaction. Establish centralized registries (Hugging Face or similar) to reduce duplication and broaden access. The goal is shared infrastructure, not proprietary advantage.

Develop community norms for ethical capture. We propose guidelines in Appendix D addressing consent frameworks, legal considerations, and privacy protections. The community should refine these norms through open discussion, balancing research value against potential harms.

6. Conclusion

The web agent research community has made real progress, but handcrafted replicas are now the bottleneck. They are slow and expensive to build, lock us into small task sets, and create a two-tier ecosystem where only well-funded organizations can afford realistic environments.

Capture-replay offers a path forward: fast collection from real work, economic grounding by construction, and open sharing to democratize access. The remaining challenges are about collection scale and responsible sharing, not environment engineering.

Our position, restated: The web agent research community should prioritize capture-replay infrastructure over handcrafted replicas. Shift investment toward capture tooling and large-scale collection, and standardize sharing so long-horizon, economically valuable tasks become accessible to everyone. Capture-replay has real limitations (human trajectories cannot anticipate every valid agent strategy), but it makes coverage gaps measurable and addressable. The tools exist; the decision is whether we continue paying the replica tax or move the field forward.

Acknowledgements

[Redacted for anonymous review.]

References

- AGI Inc. <https://agi.inc>, 2025.
- Kaizen. <https://www.kaizenautomation.com>, 2025.
- Mechanize. <https://www.mechanize.work>, 2025.
- Plato. <https://plato.so>, 2025.
- Azevedo, P. T. Nock: HTTP server mocking and expectations library for Node.js. <https://github.com/nock/nock>, 2011.
- Binié, M. VCR: Record your test suite’s HTTP interactions and replay them during future test runs. <https://github.com/vcr/vcr>, 2011.
- Drouin, A. et al. WorkArena: How capable are web agents at solving common knowledge work tasks? *arXiv preprint arXiv:2403.07718*, 2024.
- Garg, D. et al. REAL: Benchmarking autonomous agents on deterministic simulations of real websites. *arXiv preprint arXiv:2504.11543*, 2025.
- He, H. et al. WebVoyager: Building an end-to-end web agent with large multimodal models. *arXiv preprint arXiv:2401.13919*, 2024.
- Li, Z. et al. Mind2Web 2: Evaluating agentic search with agent-as-a-judge. *arXiv preprint arXiv:2506.21506*, 2025.
- METR. Measuring AI ability to complete long tasks. <https://metr.org/blog/2025-03-19-measuring-ai-ability-to-complete-long-tasks/>, 2025.
- Microsoft. Playwright: Mock APIs: HAR replay. <https://playwright.dev/docs/mock#mocking-with-har-files>, 2023.
- Netflix. Polly.JS: Record, replay, and stub HTTP interactions. <https://netflix.github.io/pollyjs/>, 2018.
- O’Callahan, R., Jones, C., Froyd, N., Huey, K., Noll, A., and Partush, N. Lightweight user-space record and replay. In *Proceedings of the 2017 USENIX Annual Technical Conference (ATC)*, pp. 551–564, 2017.
- Song, Y. et al. BEARCUBS: A benchmark for computer-using web agents. *arXiv preprint arXiv:2503.07919*, 2025.
- Wei, J. et al. BrowseComp: A simple yet challenging benchmark for browsing agents. *arXiv preprint arXiv:2504.12516*, 2025.
- Xie, T. et al. OSWorld: Benchmarking multimodal agents for open-ended tasks in real computer environments. *arXiv preprint arXiv:2404.07972*, 2024.
- Xu, F. F. et al. TheAgentCompany: Benchmarking LLM agents on consequential real world tasks. *arXiv preprint arXiv:2412.14161*, 2024.
- Xue, T. et al. An illusion of progress? assessing the current state of web agents. *arXiv preprint arXiv:2504.01382*, 2025.
- Zeff, M. Silicon valley bets big on ‘environments’ to train AI agents. *TechCrunch*, September 2025. <https://techcrunch.com/2025/09/21/silicon-valley-bets-big-on-environments-to-train-ai-agents/>
- Zhang, C. et al. Mind2Web: Toward a generalist agent for the web. *arXiv preprint arXiv:2306.06070*, 2023.
- Zhou, S. et al. WebArena: A realistic web environment for building autonomous agents. *arXiv preprint arXiv:2307.13854*, 2023.

A. TRACE Implementation Details

This appendix provides technical details on the TRACE implementation for researchers interested in understanding, extending, or reproducing the system. The complete source code is available at the anonymous repository linked in Section 3.




Figure 1. TRACE pipeline: collection, post-processing, offline replay, and agent evaluation.

A.1. Collection Architecture

TRACE’s collection system is built on Playwright with a stealth-configured Chromium browser designed to avoid anti-automation detection while maintaining full recording fidelity.

Browser Configuration. The collector launches Chromium with carefully selected arguments to evade bot detection:

```
--disable-blink-features=
AutomationControlled
--disable-features=VizDisplayCompositor
--no-sandbox
--disable-web-security
--enable-automation=false
```

The context is configured with realistic defaults: 1366×768 viewport, en-US locale, America/NewYork timezone, and standard HTTP headers including Accept-Language, Sec-Fetch-* headers, and geolocation permissions.

Event Recording. The Recorder class captures events across multiple categories:

For each triggering event, the recorder captures:

1. **Accessibility Snapshot:** A YAML-formatted tree of up to 400 interactive elements extracted via Playwright’s accessibility API, including element roles, names, ARIA attributes, and unique reference IDs.
2. **Screenshots:** Captured via Chrome DevTools Protocol (CDP) to avoid visual flicker, throttled to maximum one per 500ms to prevent duplicates.

3. **Network Traffic:** Full HAR (HTTP Archive) recording of all requests and responses, including headers, POST data, and response bodies (base64-encoded for binary content).
4. **Browser State:** Cookies, localStorage, sessionStorage, and IndexedDB snapshots captured at task completion.




Figure 2. Collecting a task using the Tkinter desktop app.

Data Storage. Each collection produces:

- SQLite database (`tasks.db`) with task metadata and step records
- Per-step DOM snapshots in `data/doms/task-{id}/step-{n}.txt`
- Screenshots in `data/screenshots/task-{id}/*.png`
- Video recordings in `data/videos/task-{id}/*.webm`
- Capture bundle in `data/captures/task-{id}/` containing: `manifest.json` (task metadata and resource index), `recording.har` (complete HTTP archive), `storage/storage.state.json` (browser storage snapshot), and `resources/` (cached static assets)

A.2. Post-Processing Pipeline

The post-processing pipeline transforms raw captures into clean, shareable artifacts through four sequential stages.

Stage 1: Tool-Call Parsing. Raw browser events are converted to a standardized Domain-Specific Language (DSL):

The parser handles event accumulation (e.g., multiple key-down events into a single `type` action), detects navigation consequences of clicks, and preserves DOM snapshot references for each action.

Table 4. Event categories captured by TRACE.

Category	Event Types	Trigger
State:Page	load, domcontentloaded, loaded	Page lifecycle events
State:Browser	navigated, navigate_start, back	Navigation events
Action:User	click, input, contextmenu, submit	User interactions

Table 5. Tool-call DSL mapping from browser events.

DSL Action	Source Events	Parameters
go_to	state:browser:navigated (initial)	url
click	action:user:click, pointerdown/up	coordinates, element_info, navigates_to
type	action:user:input, keydown	value, submit, element_info
scroll	action:user:scroll	x, y (absolute coordinates)

Stage 2: Credential Extraction. An LLM-based extractor identifies credential-related interactions by analyzing:

- Tool calls with `type` actions on login-related elements
- DOM context around input fields (labels, placeholders, ARIA descriptions)
- Navigation patterns indicating authentication flows

Extracted credentials are stored in a structured format:

```
{
  "website": "amazon.com",
  "fields": [
    { "field_name": "email", "field_value": "user@example.com" },
    { "field_name": "password", "field_value": "[REDACTED]" }
  ],
  "tool_call_ids": [3, 4]
}
```

This separation enables credential substitution during sharing while preserving the trajectory structure.

Stage 3: Checkpoint Selection. An LLM analyzes the full trajectory to identify 2 to 3 semantically meaningful intermediate states suitable for partial-credit evaluation. The prompt instructs the model to consider:

- Key actions that represent partial task completion
- Navigation events to significant pages (results, confirmation screens)
- Timestamp gaps indicating complex reasoning or significant progress

Stage 4: Ignore-List Construction. Network traffic is analyzed to identify hosts and URL patterns that can be safely omitted during replay:

- Analytics and tracking domains (Google Analytics, Facebook Pixel, etc.)
- Ad networks and retargeting services
- Non-essential third-party services (chat widgets, A/B testing)

In practice, this stage discards approximately 50% of captured HTTP requests. Modern websites include substantial tracking and analytics traffic that is entirely unnecessary for core functionality. The resulting `ignored.json` file lists URL patterns to abort during replay, significantly reducing noise and improving determinism.

A.3. Replay Engine

The replay module reconstructs captured sessions for offline evaluation through a multi-stage request matching system.



Figure 3. Offline replay and tool-call parsing for an Amazon task.

Request Routing. When the browser issues a request during replay:

1. **Ignore Check:** URLs matching patterns in `ignored.json` are immediately aborted.
2. **Exact Match:** The HAR file is searched for entries with identical method and URL base (scheme + host + path). If a single match exists, it is used directly.

3. **Character-Based Matching:** For URLs with dynamic query parameters, a character-frequency similarity score is computed:

$$\text{score} = \sum_c \min(\text{tgt}[c], \text{cand}[c])$$

where $\text{tgt}[c]$ and $\text{cand}[c]$ are the character counts for character c in the target and candidate URLs respectively. URLs are normalized before matching (removing timestamp parameters, sorting query strings). Candidates with greater than 90% character overlap and matching all target characters are treated as perfect matches.

4. **LLM Disambiguation:** When multiple candidates remain after character-based filtering, the top-5 candidates (ranked by match score) are sent to an LLM for selection. The prompt provides: target request details (method, normalized URL, headers, POST data), candidate request details with response MIME types, and character match scores as additional context.

Response Fulfillment. Once a HAR entry is selected: (1) response headers are extracted, with `Set-Cookie` headers combined with newlines (per HTTP spec); (2) response bodies are decoded (base64 for binary content, UTF-8 for text); (3) the request is fulfilled via Playwright’s route API with appropriate content-type headers.

Caching. LLM match decisions are cached in `matches.json` keyed by `{method}-{url}-{body_hash_16}`. Subsequent replays reuse cached decisions for determinism. The `--ignore-cache` flag forces fresh LLM matching for debugging.

Storage State. Optional `--include-storage-state` flag restores cookies and `localStorage` from the capture, enabling replay of authenticated sessions.

A.4. Evaluation Harness

TRACE includes a minimal evaluation runner demonstrating integration with browser-use agents:

1. **Environment Launch:** Load capture bundle, configure HAR replay routing, open starting URL.
2. **Agent Execution:** Initialize browser-use agent with task description and target website. Agent interacts with the replayed environment through standard browser automation APIs.

3. **Checkpoint Evaluation:** Compare agent’s trajectory against human checkpoints using semantic similarity. An LLM judge determines whether the agent reached equivalent states, allowing alternative valid paths.

4. **Success Determination:** Binary success is assessed by comparing the agent’s final state against the task’s expected outcome (for action tasks) or answer (for information retrieval tasks).

B. Comparison with HTTP Record-Replay Tools

Table 6 compares TRACE’s replay approach with existing HTTP record-replay tools. While these tools share the fundamental capture-replay paradigm, they differ significantly in request matching sophistication and intended use case.

Key differences. **Polly.js** (Netflix, 2018) uses adapters (fetch, XHR, Puppeteer) to intercept HTTP traffic and records to HAR format. Matching is strict by default: requests must match method, URL, headers, and body exactly. Dynamic parameters require manual configuration of “matchRequestsBy” rules to ignore specific fields. Designed for deterministic test suites where request sequences are predictable.

VCR (Binié, 2011) pioneered the “cassette” metaphor for recording HTTP interactions in Ruby. It offers configurable matchers (method, URI, host, path, body, headers) and supports regex patterns for dynamic content. However, matching remains deterministic; there is no semantic understanding of request equivalence. The “record: :new_episodes” mode can extend cassettes but does not handle request re-ordering.

Nock (Azevedo, 2011) provides HTTP mocking for Node.js with interceptors that match URL patterns (including regex) and request bodies. By default, interceptors are consumed after use (one-shot), which aligns with TRACE’s index tracking but requires knowing the exact request sequence in advance. No fallback mechanism exists when strict matching fails.

Playwright HAR (Microsoft, 2023) offers “route-FromHAR” for replaying captured traffic. The “url” option supports glob patterns (e.g., “**/api/**”) but not semantic matching. The “update” mode re-records missing requests to live servers, which is useful for extending captures but not for offline-only evaluation. No mechanism exists for disambiguating between multiple matching entries.

TRACE addresses the gap between these unit-test-oriented tools and the requirements of web agent evaluation. Agents explore unpredictably, generating request orderings that differ from the recorded trajectory. TRACE’s character-

Table 6. Comparison of HTTP record-replay tools. TRACE adds semantic matching capabilities designed for the non-determinism inherent in web agent evaluation.

Tool	Matching Strategy	Dynamic Params	Reuse Prevention	Primary Use Case
Polly.js	Strict URL + method	Manual config	No	JS unit testing
VCR	Configurable matchers	Regex patterns	Optional	Ruby integration tests
Nock	URL + method + body	Regex/functions	One-shot default	Node.js mocking
Playwright HAR	Strict or glob pattern	Basic wildcards	No	E2E test mocking
TRACE	Multi-tier + LLM	Auto char-similarity	Yes (index tracking)	Web agent evaluation

frequency matching handles URL variations from timestamps and session tokens without manual configuration. LLM disambiguation resolves ambiguous cases where multiple HAR entries could satisfy a request. Index tracking ensures that requests to the same URL at different points in a session receive appropriate responses, which is critical for authentication flows where a URL may be requested before and after login with different expected responses.

C. Captured Environment Statistics

This appendix provides detailed statistics on the six demonstration environments released with TRACE. All environments are available at the anonymous dataset repository linked in Section 3. Statistics are reproduced from our proof-of-concept dataset collection.

Column definitions:

- **Clicks:** Number of click events recorded
- **Tools:** Number of high-level tool calls after post-processing
- **Creds:** Whether the task involved credential entry (login flows)
- **HTTP:** Total HTTP request/response pairs in HAR file (before ignore-list filtering)
- **Cache:** Number of cached static resources
- **Shots:** Number of screenshots captured
- **DOM:** Number of DOM/accessibility snapshots captured
- **Time(s):** Task execution duration in seconds (successful run only)
- **Size:** Total capture bundle size on disk

Aggregate totals: Total HTTP requests: 7,725 (before filtering; ~50% discarded by ignore-list); Total cached resources: 3,092; Total screenshots: 120; Total DOM snapshots: 160; Total storage: 663 MB.

Note on collection time: The Time(s) column reflects the duration of a successful task execution. Actual collection effort is approximately 6:40 per task when accounting for retries due to collection errors, website issues, or mistakes during demonstration. The six-task dataset required about 40 minutes of total collection time.

C.1. Task Descriptions

Task 1: GitHub (Authenticated). Sign in, star a repository, search for and follow a user. Exercises authentication, navigation, and account actions.

Task 2: Amazon (Authenticated). Sign in, navigate to deals, find discounted item, add to cart, proceed toward checkout. Exercises e-commerce workflow including authentication, filtering, and cart management.

Task 3: Ultimate Guitar (Authenticated). Sign in, search for tabs, interact with playback controls. Exercises authentication on media-heavy site.

Task 4: Uniqlo (Unauthenticated). Browse catalog, apply filters, view product details, add to cart. Exercises e-commerce browsing without authentication.

Task 5: Kayak (Information Retrieval). Search flights with constraints, apply filters, compare results. Exercises travel search with complex filtering.

Task 6: Airbnb (Information Retrieval). Search accommodations with filters, browse listings with map interaction. Exercises map-based search with filtering.

C.2. Observations

Several patterns emerge from the collected data:

1. **HTTP traffic scales with site complexity.** Ultimate Guitar generated the most HTTP requests (2,752), likely due to media-heavy content and third-party integrations. Travel sites (Kayak, Airbnb) had moderate traffic despite complex UIs.
2. **Bundle size correlates with cached resources.** Uniqlo and Kayak, despite different task types, both required substantial cached resources (783 and 452

Table 7. Complete statistics for the six captured demonstration environments.

Task	Website	Type	Clicks	Tools	Creds	HTTP	Cache	Shots	DOM	Time(s)	Size
1	github.com	Action	11	8	Yes	715	311	12	35	39.0	64 MB
2	amazon.com	Action	14	11	Yes	1,199	468	20	31	69.2	94 MB
3	ultimate-guitar.com	Action	19	15	Yes	2,752	534	20	46	54.9	117 MB
4	uniqlo.com	Action	11	10	No	1,315	783	12	12	46.1	128 MB
5	kayak.com	Info	14	11	No	816	452	32	17	110.7	154 MB
6	airbnb.com	Info	15	12	No	928	544	24	19	103.9	106 MB

respectively), resulting in larger bundle sizes.

- 3. Authenticated tasks were faster.** Tasks 1 to 3 (authenticated) averaged 54.4 seconds, while tasks 5 to 6 (information retrieval, unauthenticated) averaged 107.3 seconds. This likely reflects the more exploratory nature of information retrieval tasks.
- 4. DOM snapshot frequency varies by interaction pattern.** Ultimate Guitar (46 snapshots) and GitHub (35 snapshots) had higher snapshot counts due to more frequent triggering events (clicks, form submissions).

Even accounting for retries and mistakes, total collection effort for all six environments was approximately 40 minutes (~6:40 per task), demonstrating the efficiency of capture-replay compared to handcrafted environment construction which can require hours or days per task.

C.3. Replay Validation Results

To validate replay correctness, we manually traversed each captured environment following the original trajectory. Table 8 reports matching outcomes across the multi-tier system.

Observations. Matching success varies substantially by site complexity. Simple authenticated flows (GitHub, Amazon) achieve near-perfect deterministic matching (99%+). Sites with heavy dynamic content (Ultimate Guitar’s ad network, Uniqlo’s tracking pixels) require more LLM disambiguation but still resolve correctly.

Temporal mismatch. Airbnb’s high LLM requirement (25.3%) stems from temporal mismatch: the task involved date-dependent search, and replay occurred weeks after capture. Airbnb encodes search dates directly in API URLs, causing requests to diverge from recorded entries. Kayak, despite similar datepicker interaction, encodes dates differently and achieved 89.6% deterministic matching. This illustrates the temporal validity limitation discussed in Section 4: date-sensitive tasks degrade faster than static workflows.

Failed requests. The 1 to 2% failure rate on three tasks reflects requests for which no plausible HAR entry existed,

typically dynamically-generated tracking pixels or ad callbacks absent from the original capture. These failures did not prevent task completion; the affected resources were non-essential (analytics, ads). Production systems should extend the ignore-list to cover such cases.

D. Ethical and Safety Guidelines

This appendix proposes guidelines for responsible capture and distribution of browser environments.

D.1. Legal Considerations

Copyright. Captured environments contain copyrighted material. Fair use doctrines may permit research use, but this is untested for AI training. We recommend documenting fair use rationale and limiting distribution to research contexts. Notably, replica-based benchmarks face identical concerns: WebArena reproduces Reddit’s interface; REAL simulates commercial sites.

Terms of service. Most websites prohibit automated access in ToS, though enforceability varies (*hiQ v. LinkedIn*). We recommend avoiding high-volume systematic capture and respecting robots.txt.

CFAA considerations. Post-Van Buren (2021), mere ToS violations are not CFAA crimes, but circumventing access controls may be. Capture only from accounts you own or have authorization to use.

D.2. Privacy and Data Protection

Captures contain personal data requiring attention under GDPR/CCPA:

- 1. Informed consent.** Demonstrators should consent to capture and distribution.
- 2. PII detection and redaction.** Post-processing should detect and redact names, emails, addresses, financial data, and government identifiers.
- 3. Third-party minimization.** Avoid capturing pages with extensive third-party PII unless necessary.

Table 8. Replay matching breakdown for human traversal of captured environments. Deterministic matches require no LLM disambiguation; LLM-required matches were resolved correctly; failed matches had no suitable HAR entry.

Task	Website	Requests	Deterministic	LLM Required	Failed	LLM Calls	Avg Confidence
1	GitHub	725	99.9%	0.1%	0%	1	0.68
2	Amazon	676	99.1%	0.9%	0%	6	0.48
3	Ultimate Guitar	342	88.6%	11.4%	2.0%	39	0.59
4	Uniqlo	663	94.1%	5.9%	1.2%	39	0.55
5	Kayak	211	89.6%	10.4%	0%	22	0.57
6	Airbnb	501	74.7%	25.3%	1.8%	127	0.55
Total/Avg		3,118	91.3%	7.5%	1.1%	234	0.55

4. **Credential extraction.** TRACE separates authentication data from content; extend this to general PII.

PCI-DSS), legal review. Distribution: institutional agreements.

D.3. Security Risks

Credential exposure. Captures may contain passwords, session cookies, API keys, and tokens. Mitigations: mandatory credential extraction, secure storage, credential rotation after capture, and preference for short-lived tokens.

Vulnerability disclosure. Captures may reveal security flaws. Review captures before distribution; establish responsible disclosure procedures; withhold captures revealing serious vulnerabilities.

D.4. Misuse Potential

Capture-replay has dual-use risks. Environments could train agents for automated fraud, credential stuffing, phishing, or surveillance.

Mitigations: Access controls for sensitive domains (finance, healthcare, government); licensing that prohibits malicious applications; consider whether certain environment types (payment completion, account creation) should be captured at all; tiered access balances openness with responsibility.

Our view: Proprietary concentration does not eliminate misuse risk; it prevents independent safety research. Transparent infrastructure enables collective norm development.

D.5. Three-Tier Consent Framework

1. **Tier 1: Self-demo.** Scope: own accounts. Requirements: consent, credential extraction. Distribution: public after processing.
2. **Tier 2: Authorized.** Scope: others' demos, research accounts. Requirements: IRB review, written consent, PII redaction. Distribution: research-use license.
3. **Tier 3: Sensitive.** Scope: healthcare, finance, government. Requirements: domain compliance (HIPAA,

D.6. Domain-Specific Notes

- **Healthcare:** Never capture real patient data; treat as Tier 3
- **Financial:** Stop before payment completion; never capture real card data
- **Government:** Jurisdiction-specific regulations apply
- **Enterprise SaaS:** Obtain written authorization from system owners

D.7. Community Infrastructure

We call for: shared registries with access controls (e.g., Hugging Face), open-source compliance tooling (PII detection, credential extraction), community review for sensitive domains, and incident response procedures.

D.8. Limitations

These guidelines are not legal advice, cannot anticipate all risks, and depend on community adoption. Strict controls may recreate capability concentration. Despite limitations, explicit guidelines are preferable to the implicit norms around replicas, which sidestep these questions entirely.